

Project 3

Jake Carlson

November 4, 2017

Abstract

This report examines federal payroll data in the years 2005 and 2013. The data will be transformed into transaction data so that association rule mining can be performed. By looking at the differences in maximal, closed, and frequent itemsets, I will be able to understand common features about employees under the presidencies of George W. Bush and Barack Obama.

Contents

1	Business Understanding	3
2	Data Preparation	3
3	Modeling	6
3.1	Frequent Itemsets	6
4	Evaluation	9

1 Business Understanding

In this report I will again be examining the federal payroll data obtained by BuzzFeed News through the Freedom of Information Act. I will look at the years 2005 and 2013 so I can gauge how the government changed as the presidency transitioned from George W. Bush to Barack Obama. This report will focus on association rule mining. The payroll data will be formatted as transaction data, where each employee is treated as a bag of features where all attributes that apply to them are placed in their bag. With the data formatted in this way, I can use the `arules` package to find items that occur together. These items can be treated as rules, where the occurrence of some items tend to imply the occurrence of another item. With these rules, along with statistics about the rule such as the support, or the ratio of instances that express the rule, and confidence, or the conditional probability that the predicted item occurs in an instance which contains all remaining items for the rule, I will build a picture of the key differences between groups of federal employees under the two presidents.

2 Data Preparation

I will start with the data as I prepared it for classification in Project 2. This data has been cleaned such that all unknown values were replaced with NA, Age and Length of Service were adjusted to take the middle year for the range of years given in the raw data, and Pay has been discretized into the pay ranges given in Table 1. In addition, Education is discretized into the groups given in Table 2.

I will get rid of attributes that I am not interested in for association rule mining. I will drop the Agency attribute because AgencyName already encodes this data. Likewise, I will drop Station because region records the state an employee works in.

All numeric fields will be discretized based on frequency ***explain why***. SupervisoryStatus must be treated as a factor so that general employees are not grouped with supervisors. The final list of attributes is given in Table 3.

These attributes are then transformed into a transaction item matrix. The most frequent items show us what attributes are the most common for employees to have. Summaries for both the 2005 and 2013 transaction data sets are given in Table 4. We see that for both years, the most common employees are non-supervisors. The majority of employees are also college educated and have been working in the federal government for less than twelve years. Most employees are younger than 47 years old and the most frequent employment category is Administrative.

Pay Ranges
<30k
30-50k
50-70k
70-90k
90-110k
>110k

Table 1: The Pay Ranges Used To Descretize Pay

Group	Education Levels	Description
Elm	0, 1	Reached or completed elementary school
HS	3, 4, 5, 6	Reached or completed high school or an occupational program
Col	7, 8, 9, 10, 11, 12, 13	Reached or completed college with a Bachelor's degree
Grad	14, 15, 16, 17, 18, 19, 20	Any level of graduate studies, excluding a Doctorate
Doc	21, 22	A Doctorate or Post-Doctorate degree

Table 2: Ordinal Education Groups

Attribute	Scale	Range
AgencyName	Nominal	The name of each agency
region	Nominal	The name of the state
Age	Interval	[17,47), [47,57), [57,75]
Education	Ordinal	Elm, HS, Col, Grad, Doc
LOS	Interval	[1,22), 22, [27,35]
Category	Nominal	P, A, T, C, O, B
Pay	Ordinal	<30k, 30-50k, 50-70k, 70-90k, 90-110k, >110k
SupervisoryStatus	Nominal	2, 4, 5, 6, 7, 8

Table 3: Final Data Set Attributes

2005	Elements	4,720,680
	Item	Count
	SupervisoryStatus=8	4,081,468
	Education=Col	2,190,442
	Age=[17,47)	2,087,004
	LOS=[1,12)	1,900,001
	Category=A	1,716,162
	Itemset Length	Count
	4	18
	5	2529
2013	6	85,052
	7	2,034,281
	8	2,598,800
	Elements	5,323,899
	Item	Count
	SupervisoryStatus=8	4,502,511
	Education=Col	2,574,899
	LOS=[1,12)	2,568,556
	Age=[17,47)	2,276,609
	Category=A	2,034,255
	Itemset Length	Count
	5	126
	6	13,770
	7	2,173,306
	8	3,136,697

Table 4: Summaries of 2005 and 2013 Transaction Data Sets

3 Modeling

In order to mine the transaction data for meaningful rules, I will examine frequent, closed, and maximal itemsets.

3.1 Frequent Itemsets

The most frequent itemsets show the attributes that are shared by the most number of employees. The most frequent itemsets have the highest support value out of all of the itemsets. I will look for itemsets that occur for 1% of all employees. It is also useful to view itemsets with more items to see what attributes occur frequently together. The distribution of itemset sizes is given in Figure 1. We see that the distribution of sizes is consistent between the two years.

In order to gauge differences between the two administrations, I will examine the most frequent itemsets which contain more than five items. The relative item frequency for items in itemsets with length six are given in Figure 2 for 2005 and Figure 3 for 2013.

A number of trends are apparent in these graphs. We see an increase in the ratio of employees who have a Graduate level education, as well as a reduction in the ratio of employees who have a High School level education under the Obama administration. Also, there is a decrease in the ratio of employees who have a Length of Service between 12 and 22 years, with an increase in employees who have worked for the government for less than 12 years. This is consistent with Project 1 where I found that Obama employed newer employees, while employees hired before the Bush administration looked for employment elsewhere [1].

Another noticable trend is the frequency of employees who worked for the Transportation Security Administration. Under Bush, the number of employees in the TSA is quite large, reflecting the growth of the agency in response to the September 11th attacks. The TSA shrunk in size relative to the VHA under Obama. This reflects the growth of the Veterans Affairs Administration under the Obama administration.

To look closer at the relationships between items in these itemsets, I will pull out itemsets that do not reference a particular administration to examine features that are common across the entire federal government. Table 5 shows these itemsets and their support. These itemsets reflect my findings in Project 2 [2]. Notable relationships are that a higher level of education is associated with a higher annual pay when all other attributes are the same. Likewise, working for the government for longer is associated with a higher net pay. Also, The Category that applies to an employee's role is associated with different pay levels. In the examples given, Administrative roles are paid more than Technical roles. These trends hold for both administrations. The main difference is that employees in Technical roles with a High School education have slightly higher pay under Obama than under Bush.

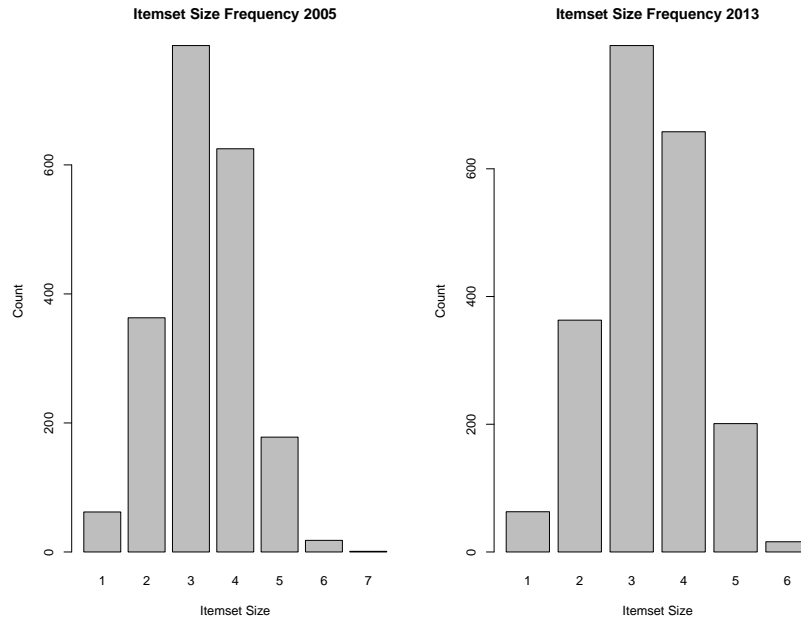


Figure 1: The Distribution of Itemset Sizes for 2005 and 2013

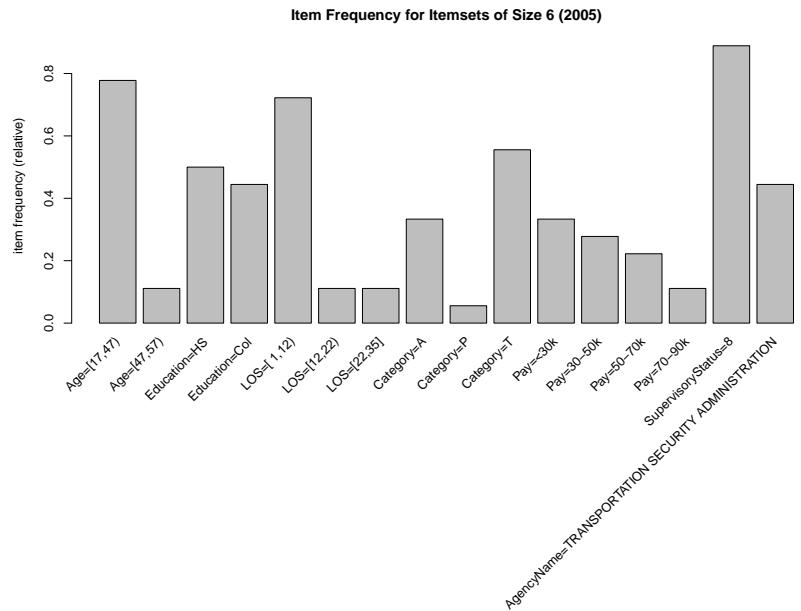


Figure 2: The Distribution of Items in Itemsets of Size Six for 2005

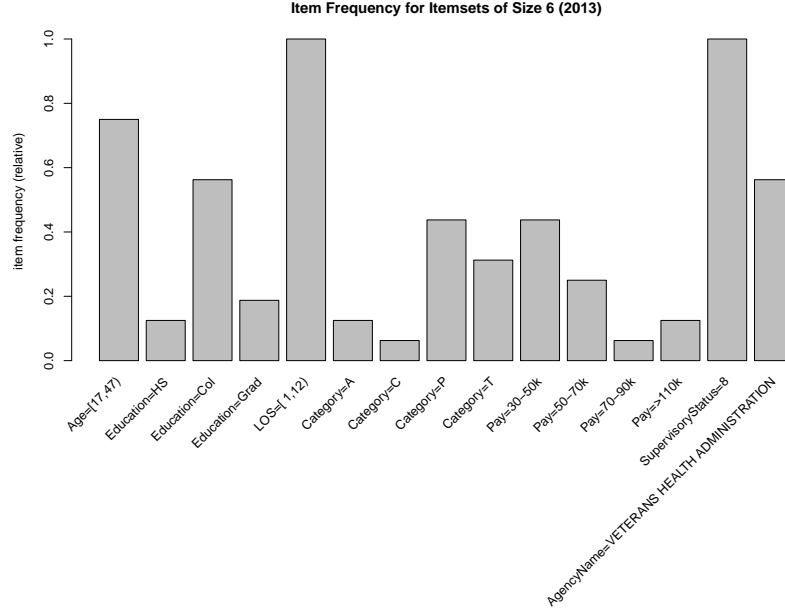


Figure 3: The Distribution of Items in Itemsets of Size Six for 2013

2005	
Itemset	Support
{Age=[17,47), Education=HS, LOS=[1,12), Category=T, Pay=30-50k, SupervisoryStatus=8}	0.0213
{Age=[17,47), Education=Col, LOS=[1,12), Category=A, Pay=50-70k, SupervisoryStatus=8}	0.0172
{Age=[17,47), Education=Col, LOS=[1,12), Category=T, Pay=30-50k, SupervisoryStatus=8}	0.0170
{Age=[17,47), Education=HS, LOS=[1,12), Category=T, Pay=<30k, SupervisoryStatus=8}	0.0165
{Age=[17,47), Education=Col, LOS=[12,22), Category=A, Pay=70-90k, SupervisoryStatus=8}	0.0133
2013	
Itemset	Support
{Age=[17,47), Education=Col, LOS=[1,12), Category=T, Pay=30-50k, SupervisoryStatus=8}	0.0234
{Age=[17,47), Education=Col, LOS=[1,12), Category=A, Pay=50-70k, SupervisoryStatus=8}	0.0192
{Age=[17,47), Education=HS, LOS=[1,12), Category=T, Pay=30-50k, SupervisoryStatus=8}	0.0178
{Age=[17,47), Education=Col, LOS=[1,12), Category=A, Pay=70-90k, SupervisoryStatus=8}	0.0165
{Age=[17,47), Education=Col, LOS=[12,22), Category=C, Pay=30-50k, SupervisoryStatus=8}	0.0133

Table 5: Frequent Itemsets of Size Six in 2005 and 2013

4 Evaluation

References

- [1] Jake Carlson *CSE 5331 - Data Mining Project 1*
<https://github.com/jakecarlson1/data-mining-projects/blob/master/project-1/report/carlson-project-1.pdf>
- [2] Jake Carlson *CSE 5331 - Data Mining Project 2*
<https://github.com/jakecarlson1/data-mining-projects/blob/master/project-2/report/carlson-project-2.pdf>